

This dataset is about the direct phone call marketing campaigns, which aim to promote term deposits among existing customers, by a Portuguese banking institution from May 2008 to November 2010.

Input variables:

- 1 - age (numeric)
- 2 - job : type of job (categorical: 'admin.','blue-collar','entrepreneur','housemaid','management','retired','self-employed','services','student','technician','unemployed','unknown')
- 3 - marital : marital status (categorical: 'divorced','married','single','unknown'; note: 'divorced' means divorced or widowed)
- 4 - education (categorical: 'basic.4y','basic.6y','basic.9y','high.school','illiterate','professional.course','university.degree','unknown')
- 5 - default: has credit in default? (categorical: 'no','yes','unknown')
- 6 - housing: has housing loan? (categorical: 'no','yes','unknown')
- 7 - loan: has personal loan? (categorical: 'no','yes','unknown') # related with the last contact of the current campaign:
- 8 - contact: contact communication type (categorical: 'cellular','telephone')
- 9 - month: last contact month of year (categorical: 'jan', 'feb', 'mar', ..., 'nov', 'dec')
- 10 - day_of_week: last contact day of the week (categorical: 'mon','tue','wed','thu','fri')
- 11 - duration: last contact duration, in seconds (numeric). Important note: this attribute highly affects the output target (e.g., if duration=0 then y='no'). Yet, the duration is not known before a call is performed. Also, after the end of the call y is obviously known. Thus, this input should only be included for benchmark purposes and should be discarded if the intention is to have a realistic predictive model. # other attributes:
- 12 - campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)
- 13 - pdays: number of days that passed by after the client was last contacted from a previous campaign (numeric; 999 means client was not previously contacted)
- 14 - previous: number of contacts performed before this campaign and for this client (numeric)
- 15 - poutcome: outcome of the previous marketing campaign (categorical: 'failure','nonexistent','success') # social and economic context attributes
- 16 - emp.var.rate: employment variation rate - quarterly indicator (numeric)
- 17 - cons.price.idx: consumer price index - monthly indicator (numeric)
- 18 - cons.conf.idx: consumer confidence index - monthly indicator (numeric)
- 19 - euribor3m: euribor 3 month rate - daily indicator (numeric)
- 20 - nr.employed: number of employees - quarterly indicator (numeric) Output variable (desired target):
- 21 - y - has the client subscribed a term deposit? (binary: 'yes','no')

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
import seaborn as sns
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)

data=pd.read_csv('bank-additional-full.csv',sep=';')

df=data.copy()

df.sample(5)
```

	age	job	marital	education	default	housing	loan	contact	mon
11357	31	technician	married	high.school	no	no	no	telephone	
33955	40	admin.	single	high.school	no	no	yes	cellular	rr
19435	49	housemaid	married	university.degree	no	yes	no	cellular	as
14732	38	self-employed	married	professional.course	no	no	no	cellular	
6726	51	unemployed	married	basic.9y	unknown	yes	no	telephone	rr

```
columns=df.columns
columns
```

```
Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
       'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
       'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
       'cons.conf.idx', 'euribor3m', 'nr.employed', 'y'],
      dtype='object')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   41188 non-null  int64
1   job                   41188 non-null  object
2   marital               41188 non-null  object
3   education             41188 non-null  object
4   default               41188 non-null  object
5   housing               41188 non-null  object
6   loan                  41188 non-null  object
7   contact               41188 non-null  object
8   month                 41188 non-null  object
9   day_of_week           41188 non-null  object
10  duration              41188 non-null  int64
11  campaign              41188 non-null  int64
12  pdays                 41188 non-null  int64
13  previous              41188 non-null  int64
14  poutcome              41188 non-null  object
15  emp.var.rate          41188 non-null  float64
16  cons.price.idx         41188 non-null  float64
17  cons.conf.idx         41188 non-null  float64
18  euribor3m             41188 non-null  float64
19  nr.employed           41188 non-null  float64
20  y                     41188 non-null  object
dtypes: float64(5), int64(5), object(11)
memory usage: 6.6+ MB
```

```
int_cols=[]
obj_cols=[]
flo_cols=[]
for column in df.columns:
    if df[column].dtype == 'int64':
        int_cols.append(column)
    elif df[column].dtype=='object':
        obj_cols.append(column)
    elif df[column].dtype=='float64':
        flo_cols.append(column)
```

```
for i in columns:
    unique_values = df[i].unique()
    if len(unique_values) > 30:
        print()
        print(f"{i}: {len(unique_values)}")
    else:
        print()
        print(f"{i}: {unique_values}")
```

```
age: 78
```

```

job: ['housemaid' 'services' 'admin.' 'blue-collar' 'technician' 'retired'
      'management' 'unemployed' 'self-employed' 'unknown' 'entrepreneur'
      'student']

marital: ['married' 'single' 'divorced' 'unknown']

education: ['basic.4y' 'high.school' 'basic.6y' 'basic.9y' 'professional.course'
            'unknown' 'university.degree' 'illiterate']

default: ['no' 'unknown' 'yes']

housing: ['no' 'yes' 'unknown']

loan: ['no' 'yes' 'unknown']

contact: ['telephone' 'cellular']

month: ['may' 'jun' 'jul' 'aug' 'oct' 'nov' 'dec' 'mar' 'apr' 'sep']

day_of_week: ['mon' 'tue' 'wed' 'thu' 'fri']

duration: 1544

campaign: 42

pdays: [999  6  4  3  5  1  0 10  7  8  9 11  2 12 13 14 15 16
         21 17 18 22 25 26 19 27 20]

previous: [0 1 2 3 4 5 6 7]

poutcome: ['nonexistent' 'failure' 'success']

emp.var.rate: [ 1.1  1.4 -0.1 -0.2 -1.8 -2.9 -3.4 -3.  -1.7 -1.1]

cons.price.idx: [93.994 94.465 93.918 93.444 93.798 93.2  92.756 92.843 93.075 92.893
                 92.963 92.469 92.201 92.379 92.431 92.649 92.713 93.369 93.749 93.876
                 94.055 94.215 94.027 94.199 94.601 94.767]

cons.conf.idx: [-36.4 -41.8 -42.7 -36.1 -40.4 -42.  -45.9 -50.  -47.1 -46.2 -40.8 -33.6
                -31.4 -29.8 -26.9 -30.1 -33.  -34.8 -34.6 -40.  -39.8 -40.3 -38.3 -37.5
                -49.5 -50.8]

euribor3m: 316

nr.employed: [5191.  5228.1 5195.8 5176.3 5099.1 5076.2 5017.5 5023.5 5008.7 4991.6
              4963.6]

y: ['no' 'yes']

```

Got The Unique Values from the dataframe for each columns their is no redundant value with spelling mistake

✓ Checking For Null Values

```
df.isnull().sum()
```

```

age           0
job           0
marital       0
education     0
default       0
housing       0
loan          0
contact       0
month         0
day_of_week   0
duration      0
campaign      0
pdays        0
previous      0
poutcome      0
emp.var.rate  0
cons.price.idx 0
cons.conf.idx 0
euribor3m     0
nr.employed   0
y             0
dtype: int64

```

✓ Checking Duplicate Rows

```
index=df[df.duplicated()].index

df.drop(index,inplace=True)

df[df.duplicated()].index

Index([], dtype='int64')
```

✓ EDA

Univariate Analysis


k='Black'

```
print(int_cols,'\n',flo_cols,'\n',obj_cols)

['age', 'duration', 'campaign', 'pdays', 'previous']
['emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed']
['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'day_of_week', 'poutcome', 'y']
```


AGE

```
plt.figure(figsize=(6,3))
sns.kdeplot(df['age'], color='lime', shade=True, label='Age')
plt.axvline(x=df['age'].mean(), color='k', linestyle='--', label='Mean Age: {:.2f}'.format(df['age'].mean()))
plt.legend()
plt.show()
```



Duration


```
plt.figure(figsize=(6,3))
sns.kdeplot(df['duration'], color='lime', shade=True, label='Duration')
plt.axvline(x=df['duration'].mean(), color='k', linestyle='--', label='Mean Duration: {:.2f}'.format(df['duration'].mean()))
plt.legend()
plt.show()
```



This Data is skewed towards right this means most of the clients has been never called before

Pdays

```
plt.figure(figsize=(6,3))
sns.kdeplot(df['pdays'], color='lime', shade=True, label='pdays')
plt.axvline(x=df['pdays'].mean(), color='k', linestyle='--', label='Mean Pdays: {:.2f}'.format(df['pdays'].mean()))
plt.legend()
plt.show()
```



As I told in pervious 'duration plot' most of the clients have never been contacted

This plot also make that statement true as most of the value in this pdays column are 999


Job

```
plt.figure(figsize=(6,3))

job_counts = df['job'].value_counts()

sns.barplot(x=job_counts.index, y=job_counts.values)
plt.xticks(rotation=90)

plt.xlabel('Job')
plt.ylabel('Frequency')
plt.title('Job Frequency')
plt.show()
print(job_counts)
```



```

job
admin.      10419
blue-collar  9253
technician  6739
services    3967
management  2924
retired     1718
entrepreneur 1456
self-employed 1421
housemaid   1060
unemployed  1014
student     875
unknown     330
Name: count, dtype: int64


```

MARITAL

```

plt.figure(figsize=(5,6))
mar=df['marital'].value_counts()
plt.pie(mar,labels=mar.index,autopct='%1.1f%%');

```




EDUCATION

```

edu=df['education'].value_counts()
edu
plt.figure(figsize=(5,4))
sns.countplot(x='education',data=df,color='orange')
plt.xticks(rotation=90);

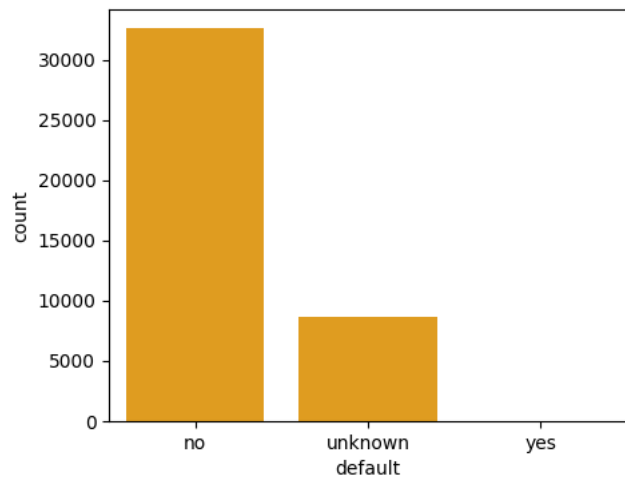
```



DEFAULT

```
plt.figure(figsize=(5,4))
sns.countplot(x='default', data=df,color='orange')
```


<Axes: xlabel='default', ylabel='count'>



HOUSING

```
housing=df['housing'].value_counts()
plt.figure(figsize=(5,4))
sns.countplot(x='housing', data=df,color='orange')
print(housing)
```


```
housing
yes      21571
no       18615
unknown   990
Name: count, dtype: int64
```



LOAN

```
plt.figure(figsize=(5,4))
sns.countplot(x='loan', data=df,color='orange')
```

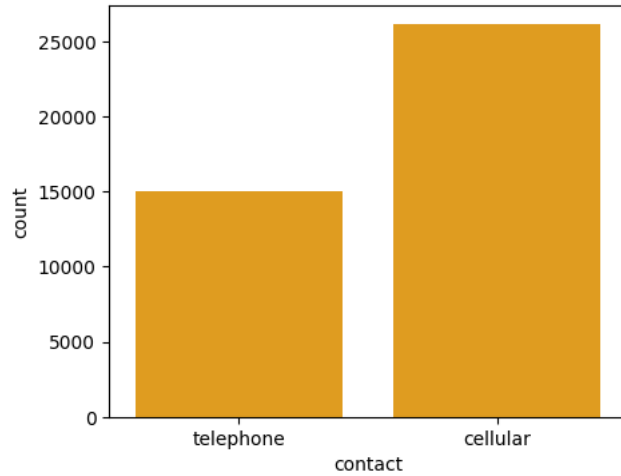
<Axes: xlabel='loan', ylabel='count'>



CONTACT

```
plt.figure(figsize=(5,4))
sns.countplot(x='contact', data=df,color='orange')
```



<Axes: xlabel='contact', ylabel='count'>



MONTH

```
x='month'
temp=df[x].value_counts()
plt.figure(figsize=(5,4))
sns.barplot(x=temp.index, y=temp.values,color='orange')
print(temp)
```

```
month
may    13767
jul     7169
aug     6176
jun     5318
nov     4100
apr     2631
oct       717
sep       570
mar       546
dec       182
Name: count, dtype: int64
```




DAY OF WEEK

```
x='day_of_week'
temp=df[x].value_counts()
plt.figure(figsize=(5,4))
sns.barplot(x=temp.index, y=temp.values,color='orange')
print(temp)
```

```

day_of_week
thu      8618
mon      8512
wed      8134
tue      8086
fri      7826
Name: count, dtype: int64

```



POUTCOME

```

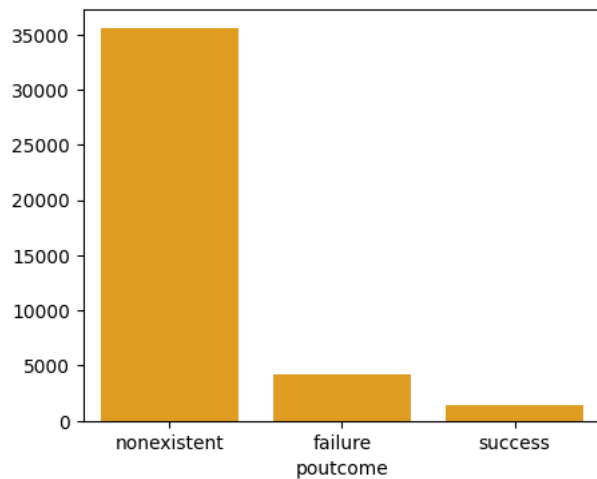
x='poutcome'
temp=df[x].value_counts()
plt.figure(figsize=(5,4))
sns.barplot(x=temp.index, y=temp.values,color='orange')
print(temp)

```

```

poutcome
nonexistent    35551
failure        4252
success        1373
Name: count, dtype: int64

```



TARGET COLUMN

```


x='y'
temp=df[x].value_counts()
plt.figure(figsize=(5,4))
sns.barplot(x=temp.index, y=temp.values,color='orange')
print(temp)

```

```

y
no    36537
yes    4639
Name: count, dtype: int64

```



More People Have rejected term loan than people accepted

Data is Uneven we will manage this by techniques like stratified cv

EMPLOYEE VARIABILITY RATE

```


x='emp.var.rate'
temp=df[x].value_counts()
plt.figure(figsize=(5,4))
sns.barplot(x=temp.index, y=temp.values,color='orange')
print(temp)

```

```


emp.var.rate
1.4    16228
-1.8    9182
1.1    7762
-0.1    3682
-2.9    1662
-3.4    1070
-1.7    773
-1.1    635
-3.0    172
-0.2     10
Name: count, dtype: int64

```




CONSUMER PRICE INDEX

```
x='cons.price.idx'
temp=df[x].value_counts()
plt.figure(figsize=(18,4))
sns.barplot(x=temp.index, y=temp.values,color='orange')
```




CONSUMER CONFIDENCE INDEX

```
x='cons.conf.idx'
temp=df[x].value_counts()
plt.figure(figsize=(18,4))
sns.barplot(x=temp.index, y=temp.values,color='orange')
```



EUROBOR

```
x='euribor3m'
temp=df[x].value_counts()
plt.figure(figsize=(18,4))
sns.scatterplot(df[x])
```



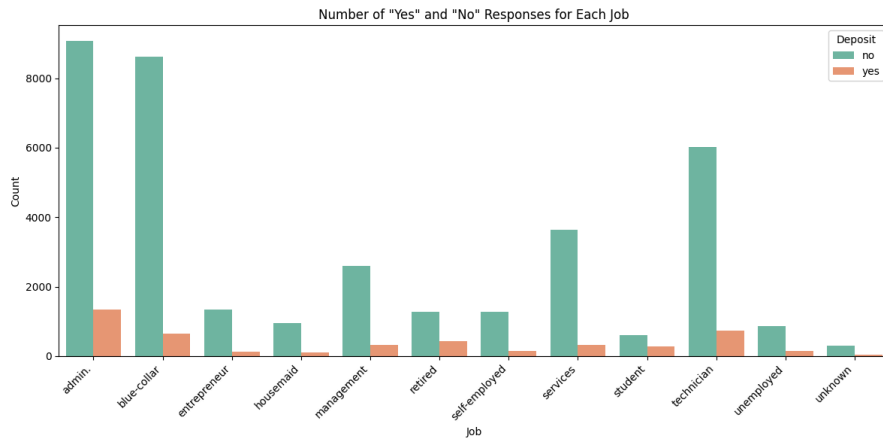
✓ BIVARIATE ANALYSIS

JOB--Y

```
# Grouping the DataFrame by 'job' and 'y', then summing the counts
grouped = df.groupby(['job', 'y']).size().unstack(fill_value=0).reset_index()

# Melt the DataFrame to have 'job' as the identifier variable
melted = pd.melt(grouped, id_vars='job', var_name='Response', value_name='Count')

# Plotting grouped bar plot
plt.figure(figsize=(12, 6))
sns.barplot(x='job', y='Count', hue='Response', data=melted, palette='Set2')
plt.title('Number of "Yes" and "No" Responses for Each Job')
plt.xlabel('Job')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right')
plt.legend(title='Deposit')
plt.tight_layout()
plt.show()
```



We can say on the basis of data that job profile admin have subscribed to term-deposit more

But Retired People, students and self-employed people are tend to subscribe more if we see their acceptance ratio

MARITAL--Y

```
gr=df.groupby(['marital', 'y']).size().unstack(fill_value=0).reset_index()
gr
```

y	marital	no	yes
0	divorced	4135	476
1	married	22390	2531
2	single	9944	1620
3	unknown	68	12

Next steps:

[Generate code with gr](#)[View recommended plots](#)

divorced,married,single,unknown WITH y

```

import matplotlib.pyplot as plt

# Define colors for the pie slices
colors = ['lightcoral', 'lightskyblue']

# Define explode for the pie slices (to emphasize 'Yes')
explode = (0, 0.1)

labels = ['No', 'Yes']

autopct = '%1.1f%'
startangle = 140

fig, axs = plt.subplots(2, 2, figsize=(5, 5))
axs = axs.flatten()


for i, (_, row) in enumerate(gr.iterrows()):
    counts = [row['no'], row['yes']]

    axs[i].pie(counts, labels=labels, colors=colors, autopct=autopct, startangle=startangle, explode=explode, shadow=True)


    axs[i].set_title(f'Marital Status: {row["marital"]}')
    axs[i].axis('equal')
plt.tight_layout()
plt.show()

```


Marital Status: divorced




Marital Status: married



Marital Status: single



Marital Status: unknown



Based On Data We can say that people with,

Marital Status =Unknown,Single Have High Subscription ratio

DEFAULT--Y

```

gr1=df.groupby(['default','y']).size().unstack(fill_value=0).reset_index()
gr1

```

y	default	no	yes
0	no	28381	4196
1	unknown	8153	443
2	yes	3	0



Next steps:

[Generate code with gr1](#)[View recommended plots](#)

```
import matplotlib.pyplot as plt

# Define colors for the pie slices
colors = ['lightcoral', 'lightskyblue']

# Define explode for the pie slices (to emphasize 'Yes')
explode = (0, 0.1)

labels = ['No', 'Yes']


autopct = '%1.1f%%'
startangle = 140

fig, axs = plt.subplots(2, 2, figsize=(5, 5))
axs = axs.flatten()

for i, (_, row) in enumerate(gr1.iterrows()):
    counts = [row['no'], row['yes']]

    axs[i].pie(counts, labels=labels, colors=colors, autopct=autopct, startangle=startangle, explode=explode, shadow=True)

    axs[i].set_title(f'Default Status: {row["default"]}')
    axs[i].axis('equal')
plt.tight_layout()
plt.show()
```



People who Have no Default are tend to subscribe term deposit more than the people who have defaulted

HOOUSING--Y

```
gr2=df.groupby(['housing', 'y']).size().unstack().reset_index()
gr2
```

y	housing	no	yes
0	no	16590	2025
1	unknown	883	107
2	yes	19064	2507



Next steps:

[Generate code with gr2](#)[View recommended plots](#)

```
import matplotlib.pyplot as plt

# Define colors for the pie slices
colors = ['lightcoral', 'lightskyblue']

# Define explode for the pie slices (to emphasize 'Yes')
explode = (0, 0.1)

labels = ['No', 'Yes']

autopct = '%1.1f%%'
startangle = 140


fig, axs = plt.subplots(2, 2, figsize=(5, 5))
axs = axs.flatten()

for i, (_, row) in enumerate(gr2.iterrows()):
    counts = [row['no'], row['yes']]

    axs[i].pie(counts, labels=labels, colors=colors, autopct=autopct, startangle=startangle, explode=explode, shadow=True)

    axs[i].set_title(f'HOUSELoan:{row["housing"]}-subs')
    axs[i].axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.

plt.tight_layout()
plt.show()
```



There is no or little relationship between housing loan and term deposit

LOAN-Y

```
gr3=df.groupby(['loan','y']).size().unstack().reset_index()
gr3
```


y	loan	no	yes
0	no	30089	3849
1	unknown	883	107
2	yes	5565	683



Next steps:

[Generate code with gr3](#)[View recommended plots](#)

```
import matplotlib.pyplot as plt

# Define colors for the pie slices
colors = ['lightcoral', 'lightskyblue']

# Define explode for the pie slices (to emphasize 'Yes')
explode = (0, 0.1)

# Define labels for the pie slices
labels = ['No', 'Yes']

autopct = '%1.1f%%'
startangle = 140


# Plotting pie charts
fig, axs = plt.subplots(2, 2, figsize=(5, 5))
axs = axs.flatten()

for i, (_, row) in enumerate(gr3.iterrows()):
    counts = [row['no'], row['yes']]

    axs[i].pie(counts, labels=labels, colors=colors, autopct=autopct, startangle=startangle, explode=explode, shadow=True)

    axs[i].set_title(f'Loan: {row["loan"]}')
    axs[i].axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.

# Adjust layout and show plot
plt.tight_layout()
plt.show()
```



No Conclusion Can be drawn as people with loan ,no loan and unknown all are subscribing term deposite with same ratio

CONTACT-Y

```
gr4=df.groupby(['contact','y']).size().unstack().reset_index()
gr4
```

y	contact	no	yes
0	cellular	22283	3852
1	telephone	14254	787



Next steps:

[Generate code with gr4](#)

[View recommended plots](#)

```
import matplotlib.pyplot as plt

# Define colors for the pie slices
colors = ['lightcoral', 'lightskyblue']

# Define explode for the pie slices (to emphasize 'Yes')
explode = (0, 0.1)

# Define labels for the pie slices
labels = ['No', 'Yes']

autopct = '%1.1f%%'
startangle = 140



fig, axs = plt.subplots(2, 1, figsize=(8, 8))
axs = axs.flatten()

for i, (_, row) in enumerate(gr4.iterrows()):
    counts = [row['no'], row['yes']]

    axs[i].pie(counts, labels=labels, colors=colors, autopct=autopct, startangle=startangle, explode=explode, shadow=True)

    axs[i].set_title(f'Contact: {row["contact"]}')
    axs[i].axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.

# Adjust layout and show plot
plt.tight_layout()
plt.show()
```



People Who Are Contacted through cellular medium are tend to subscribe for term deposit more than people contacted through telephone

MONTH-Y

```
gr5=df.groupby(['month','y']).size().unstack().reset_index()
gr5
```

y	month	no	yes
0	apr	2092	539
1	aug	5521	655
2	dec	93	89
3	jul	6521	648
4	jun	4759	559
5	mar	270	276
6	may	12881	886
7	nov	3684	416
8	oct	402	315
9	sep	314	256

Next steps:

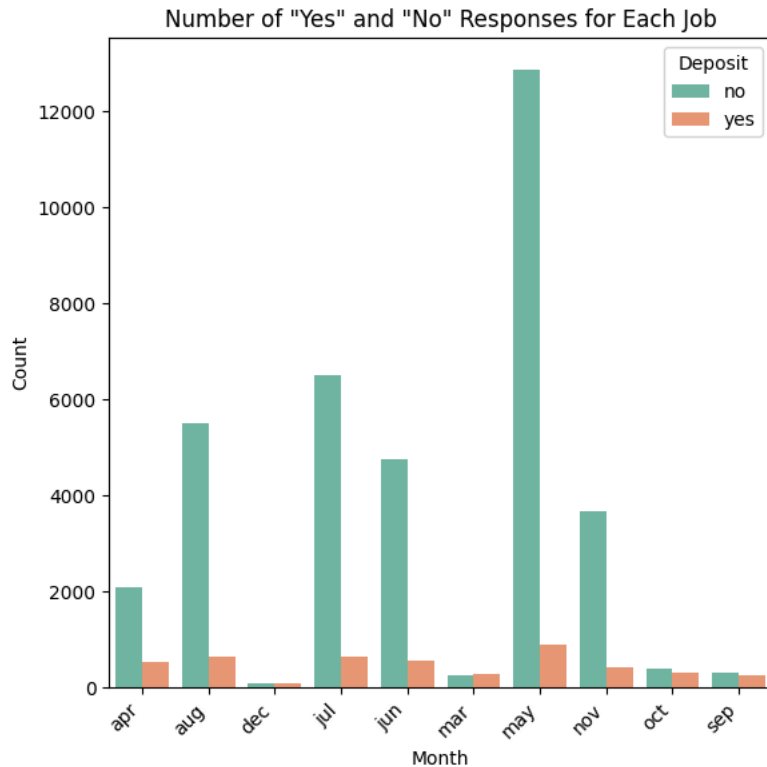
Generate code with gr5

☒ View recommended plots

```

melted = pd.melt(gr5, id_vars='month', var_name='Response', value_name='Count')
melted['percent']=round(melted['Count']*100/melted.groupby('month')['Count'].transform('sum'),1)
melted['percent']=melted['percent'].apply(lambda x: '{}%'.format(x))
# Plotting grouped bar plot
plt.figure(figsize=(6, 6))
sns.barplot(x='month', y='Count', hue='Response', data=melted, palette='Set2')
plt.title('Number of "Yes" and "No" Responses for Each Job')
plt.xlabel('Month')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right')
plt.legend(title='Deposit')
plt.tight_layout()
plt.show()

```



There is more subscription in month of may

But There is good subscription in oct,sep,april,june,july in term of no and yes ratio

```

# Grouping by 'month' and 'contact', then getting value counts
contact_counts = df.groupby(['month', 'contact']).size()

# Resetting the index to make 'month' and 'contact' as columns
contact_counts = contact_counts.reset_index(name='count')

# Summing 'count' for each month
monthly_contact_sum = contact_counts.groupby('month')['count'].sum()


new_df = pd.DataFrame(monthly_contact_sum)

new_df = new_df.rename(columns={'count': 'contacts'})

sns.barplot(x='month', y='contacts', data=new_df)

```

<Axes: xlabel='month', ylabel='contacts'>



We Can See That Number Of Contact In oct,sep,april,june,july if we increase our contact to people in these month then bank can get high subscription as in these month rejection is less as compared to other months

DAY OF WEEK--Y

```
gr6=df.groupby(['day_of_week','y']).size().unstack().reset_index()
gr6
```

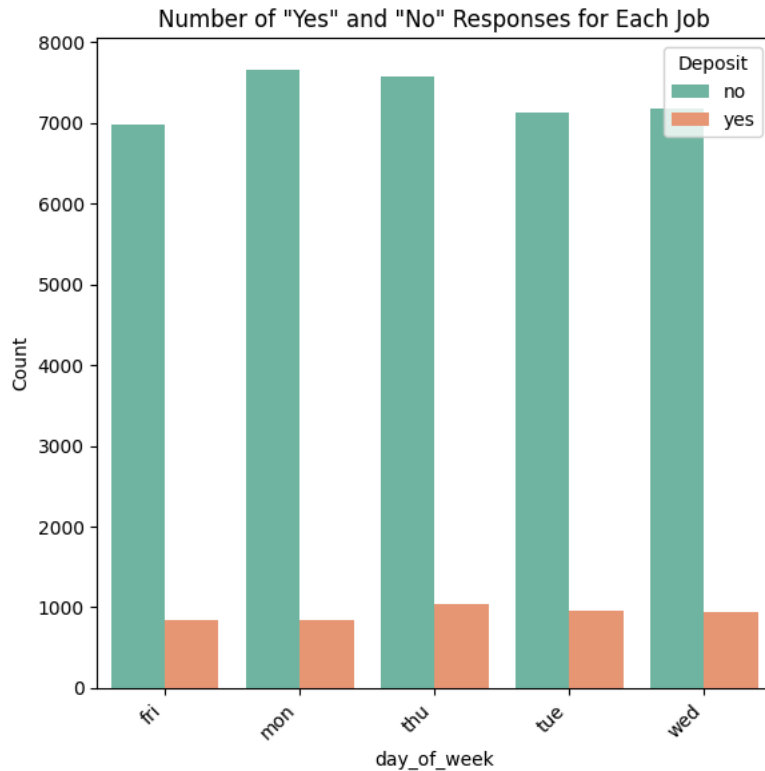
y	day_of_week	no	yes	
0	fri	6980	846	
1	mon	7665	847	
2	thu	7574	1044	
3	tue	7133	953	
4	wed	7185	949	

Next steps:

[Generate code with gr6](#)

[View recommended plots](#)

```
melted = pd.melt(gr6, id_vars='day_of_week', var_name='Response', value_name='Count')
# Plotting grouped bar plot
plt.figure(figsize=(6, 6))
sns.barplot(x='day_of_week', y='Count', hue='Response', data=melted, palette='Set2')
plt.title('Number of "Yes" and "No" Responses for Each Job')
plt.xlabel('day_of_week')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right')
plt.legend(title='Deposit')
plt.tight_layout()
plt.show()
```




People tend to Subscribe more on thursday

EMPLOYEE VARIABILITY RATE --Y

```
m=df['emp.var.rate'].mean()
above_avg=df[df['emp.var.rate']>m]
av=above_avg.groupby('y').size()
below_avg=df[df['emp.var.rate']<m]
bv=below_avg.groupby('y').size()


plt.figure(figsize=(3,3))
plt.pie(av, labels=labels, autopct='%1.1f%%', wedgeprops=dict(width=0.3), colors=colors);
plt.title('Above Avg Employee Variability---Subscription');
```

Above Avg Employee Variability---Subscription



```
plt.figure(figsize=(3,3))
plt.pie(bv, labels=labels, autopct='%1.1f%%', wedgeprops=dict(width=0.3), colors=colors)
plt.title('Below Avg Employee Variability---Subscription');
```

Below Avg Employee Variability---Subscription



When There is High Employee Variability Subscription Is Less


When There is Low Employee Variability Subscription Is More

CONSUMER PRICE INDEX---Y

```
m=df['cons.price.idx'].mean()
above_avg=df[df['cons.price.idx']>m]
av=above_avg.groupby('y').size()
below_avg=df[df['cons.price.idx']<m]
bv=below_avg.groupby('y').size()


plt.figure(figsize=(3,3))
plt.pie(bv, labels=labels, autopct='%1.1f%%', wedgeprops=dict(width=0.3), colors=colors)
plt.title('Below Avg cons.price.idx---Subscription');
```

Below Avg cons.price.idx---Subscription



```
plt.figure(figsize=(3,3))
plt.pie(av, labels=labels, autopct='%1.1f%%', wedgeprops=dict(width=0.3), colors=colors)
plt.title('Above Avg cons.price.idx---Subscription');
```

Above Avg cons.price.idx---Subscription



We Can See When Consumer Price Index is Below Avg Subscription Increases

We Can See When Consumer Price Index is Above Avg Subscription Decreases

CONSUMER CONFIDENCE INDEX --Y


```

m=df['cons.conf.idx'].mean()
above_avg=df[df['cons.conf.idx']>m]
av=above_avg.groupby('y').size()
below_avg=df[df['cons.conf.idx']<m]
bv=below_avg.groupby('y').size()

plt.figure(figsize=(3,3))
plt.pie(bv, labels=labels, autopct='%1.1f%%', wedgeprops=dict(width=0.3), colors=colors)
plt.title('Below Avg cons.conf.idx---Subscription');

```

Below Avg cons.conf.idx---Subscription




```

plt.figure(figsize=(3,3))
plt.pie(av, labels=labels, autopct='%1.1f%%', wedgeprops=dict(width=0.3), colors=colors)
plt.title('Above Avg cons.conf.idx---Subscription');

```

Above Avg cons.conf.idx---Subscription



We can see that when consumer confidence index is above avg subscription increases

We can see that when consumer confidence index is below avg subscription decreases

EURO BOR-Y


```

m=df['euribor3m'].mean()
above_avg=df[df['euribor3m']>m]
av=above_avg.groupby('y').size()
below_avg=df[df['euribor3m']<m]
bv=below_avg.groupby('y').size()

plt.figure(figsize=(3,3))
plt.pie(bv, labels=labels, autopct='%1.1f%%', wedgeprops=dict(width=0.3), colors=colors)
plt.title('Below Avg euribor3m---Subscription');


```


Below Avg euribor3m---Subscription



```
plt.figure(figsize=(3,3))
plt.pie(av, labels=labels, autopct='%1.1f%%', wedgeprops=dict(width=0.3), colors=colors)
plt.title('Above Avg euribor3m---Subscription');
```

Above Avg euribor3m---Subscription



We can see when the euribor is below avg people tends to subscribe term deposit more than when euribor is above avg..

But Why So???


Low EURIBOR rates may lead to lower borrowing costs for banks, which could incentivize them to offer more favorable terms on term deposits to attract funding. Banks may also use term deposits as a source of funding for their lending activities, offering competitive rates to attract depositors.

DURATION--Y

```
m=df['duration'].mean()
above_avg=df[df['duration']>m]
av=above_avg.groupby('y').size()
below_avg=df[df['duration']<m]
bv=below_avg.groupby('y').size()


plt.figure(figsize=(3,3))
plt.pie(bv, labels=labels, autopct='%1.1f%%', wedgeprops=dict(width=0.3), colors=colors)
plt.title('Below Avg duration---Subscription');
```

Below Avg duration---Subscription



```
plt.figure(figsize=(3,3))
plt.pie(av, labels=labels, autopct='%1.1f%%', wedgeprops=dict(width=0.3), colors=colors)
plt.title('Above Avg Duration---Subscription');
```

Above Avg Duration---Subscription



We can Clearly see that there is high correlation between duration and Y column

As Duration Increases People Subscription Increases


But, we cannot use duration column in our prediction model as we have to predict the outcome will people subscribe to term deposit if yes then only we are going to call them Hence,we will drop duration column while building predictive models

HEATMAP

```
import matplotlib.pyplot as plt
import seaborn as sns

# Select only numeric columns
numeric_df = df.select_dtypes(include=['number'])

# Plot correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.show()
```



```
print('int_cols:',int_cols)
print('obj_cols:',obj_cols)
print('flo_cols:',flo_cols)
```

```
int_cols: ['age', 'duration', 'campaign', 'pdays', 'previous']
obj_cols: ['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'day_of_week', 'poutcome', 'y']
flo_cols: ['emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed']
```

ONE HOT ENCODING THE NOMINAL CATEGORICAL COLUMNS

```
columns_categorical = ['job', 'marital',
                       'education', 'default',
                       'housing', 'loan', 'contact', 'month',
                       'day_of_week', 'poutcome']
```

```
df_2=pd.DataFrame(df)
```

```
df_2.head()
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_w
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	
1	57	services	married	high.school	unknown	no	no	telephone	may	
2	37	services	married	high.school	no	yes	no	telephone	may	
3	40	admin.	married	basic.6y	no	no	no	telephone	may	
4	56	services	married	high.school	no	no	yes	telephone	may	

```
# Apply one-hot encoding using pandas.get_dummies()
encoded_df = pd.get_dummies(df[columns_categorical],drop_first=False,dtype=int)
```

```
categorical_enc=list(encoded_df.columns)
categorical_enc
```

```
['job_admin.',
 'job_blue-collar',
 'job_entrepreneur',
 'job_housemaid',
 'job_management',
 'job_retired',
 'job_self-employed',
 'job_services',
 'job_student',
 'job_technician',
 'job_unemployed',
 'job_unknown',
 'marital_divorced',
 'marital_married',
 'marital_single',
 'marital_unknown',
 'education_basic.4y',
 'education_basic.6y',
 'education_basic.9y',
 'education_high.school',
 'education_illiterate',
 'education_professional.course',
 'education_university.degree',
 'education_unknown',
 'default_no',
 'default_unknown',
 'default_yes',
 'housing_no',
 'housing_unknown',
 'housing_yes',
 'loan_no',
 'loan_unknown',
 'loan_yes',
 'contact_cellular',
 'contact_telephone',
 'month_apr',
 'month_aug',
 'month_dec',
 'month_jul',
 'month_jun',
 'month_mar',
 'month_may',
 'month_nov',
 'month_oct',
 'month_sep',
 'day_of_week_fri',
 'day_of_week_mon',
 'day_of_week_thu',
 'day_of_week_tue',
 'day_of_week_wed',
 'poutcome_failure',
 'poutcome_nonexistent',
 'poutcome_success']
```

```
categorical=list(columns_categorical)
```

```
all_columns=list(df.columns)
```

```
numerical_columns = [col for col in all_columns if col not in categorical]
```

```
numerical_columns
```

```
['age',
 'duration',
 'campaign',
 'pdays',
 'previous',
 'emp.var.rate',
 'cons.price.idx',
 'cons.conf.idx',
 'euribor3m',
```

```
'nr.employed',
'y']
```

Removing duration as it cannot be input column and y coz it is categorical column

```
# numerical_columns.remove('duration')
```

```
numerical_columns
```

```
['age',
'duration',
'campaign',
'pdays',
'previous',
'emp.var.rate',
'cons.price.idx',
'cons.conf.idx',
'euribor3m',
'nr.employed',
'y']
```

Label Encoding Y column

```
from sklearn.preprocessing import LabelEncoder
```

```
# Assuming 'column_to_encode' is the column you want to encode
column_to_encode = 'y'
```

```
# Initialize LabelEncoder
label_encoder = LabelEncoder()
```

```
# Fit and transform the column
df[column_to_encode + '_encoded'] = label_encoder.fit_transform(df[column_to_encode])
```

```
# Now df[column_to_encode + '_encoded'] contains the encoded values
```

```
df.head(1)
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_w
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	r

```
df = pd.concat([df,encoded_df], axis = 1)
```

```
df.head()
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_w
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	
1	57	services	married	high.school	unknown		no	no	telephone	may
2	37	services	married	high.school	no	yes	no	telephone	may	
3	40	admin.	married	basic.6y	no	no	no	telephone	may	
4	56	services	married	high.school	no	no	yes	telephone	may	

```
numerical_columns.remove('y')
```

```
input_col=numerical_columns+categorical_enc
```

```
df_enc=df[input_col+['_y_encoded']]
```

```
df_enc.shape
```

```
(41176, 64)
```

WE WILL DO 2 APPROACHES AS DATA IS IMBALANCED

1)Stratified Train Valid Test Split & Stratified CV

2)UNDER SAMPLING AND THEN PERFORMING CV

```
X=df_enc.drop(columns=['y_encoded'])
y=df_enc['y_encoded'].values
```

```
print(X.shape)
print(y.shape)
```

```
(41176, 63)
(41176,)
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

```
# Then, split the test set into validation and test sets (50% each)
```

```
X_valid, X_test, y_valid, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42, stratify=y_temp)
```

```
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
```

```
X_train_scaled=scaler.fit_transform(X_train)
X_valid_scaled=scaler.transform(X_valid)
X_test_scaled=scaler.transform(X_test)
```

```
X_train=pd.DataFrame(X_train_scaled,columns=X_train.columns)
X_valid=pd.DataFrame(X_valid_scaled,columns=X_valid.columns)
X_test=pd.DataFrame(X_test_scaled,columns=X_test.columns)
```

✓ Model Selection

```

from sklearn.metrics import confusion_matrix, roc_curve, auc

def calculate_classification_metrics(y_true, y_pred, threshold):
    """
    Calculate classification metrics including accuracy, precision, recall, F1-score, and AUC.

    Parameters:
    - y_true (array-like): True labels.
    - y_pred (array-like): Predicted probabilities or scores.
    - threshold (float): Threshold value for binary classification (default=0.5).

    Returns:
    - metrics (dict): A dictionary containing classification metrics.
    """
    # Convert predicted probabilities or scores to binary predictions based on the threshold
    y_pred_binary = (y_pred >= threshold).astype(int)

    accuracy = accuracy_score(y_true, y_pred_binary)
    precision = precision_score(y_true, y_pred_binary)
    recall = recall_score(y_true, y_pred_binary)
    f1 = f1_score(y_true, y_pred_binary)
    confusion = confusion_matrix(y_true, y_pred_binary)

    # Calculate AUC
    fpr, tpr, _ = roc_curve(y_true, y_pred)
    auc_value = auc(fpr, tpr) # Use the auc function from sklearn.metrics

    metrics = {
        'Accuracy': accuracy,
        'Precision': precision,
        'Recall': recall,
        'F1-score': f1,
        'Confusion Matrix': confusion,
        'AUC': auc_value
    }

    return metrics

from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB

```

✓ Logistic Regression

```

lr=LogisticRegression(random_state = 42)
lr.fit(X_train, y_train)

```

```

▼      LogisticRegression
LogisticRegression(random_state=42)

```

```

y_train_preds = lr.predict_proba(X_train)[:,-1]
y_valid_preds = lr.predict_proba(X_valid)[:,-1]

```

```

y_train_preds

```

```

array([0.15812684, 0.4709308 , 0.01367171, ..., 0.0957597 , 0.01051456,
       0.01699796])

```

```

from sklearn.metrics import roc_curve
import matplotlib.pyplot as plt


thresholds = [0, 0.3, 0.4, 0.5]
plt.figure(figsize=(8, 6))

for threshold in thresholds:
    # Plot ROC curve for training set
    y_train_pred = (y_train_preds >= threshold).astype(int)
    fpr_train, tpr_train, _ = roc_curve(y_train, y_train_pred)
    plt.plot(fpr_train, tpr_train, lw=1, color='blue', alpha=0.5)
    plt.plot(fpr_train, tpr_train, 'o', color='blue', markersize=5) # Add filled dots
    for fpr, tpr in zip(fpr_train, tpr_train):
        plt.text(fpr, tpr, f'{threshold:.1f}', color='blue', fontsize=8, ha='right', va='bottom')

    # Plot ROC curve for validation set
    y_valid_pred = (y_valid_preds >= threshold).astype(int)
    fpr_valid, tpr_valid, _ = roc_curve(y_valid, y_valid_pred)
    plt.plot(fpr_valid, tpr_valid, lw=1, color='red', alpha=0.5)
    plt.plot(fpr_valid, tpr_valid, 'o', color='red', markersize=5) # Add filled dots
    for fpr, tpr in zip(fpr_valid, tpr_valid):
        plt.text(fpr, tpr, f'{threshold:.1f}', color='red', fontsize=8, ha='right', va='bottom')

plt.plot([0, 1], [0, 1], color='gray', lw=1, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Training and Validation Sets')
plt.legend(['Training Set', 'Validation Set'], loc="lower right")
plt.show()

```



```

from sklearn.metrics import roc_auc_score, accuracy_score, precision_score, recall_score, f1_score

```

```

print('Training:', calculate_classification_metrics(y_train, y_train_preds, threshold=0.3))
print('Validation:', calculate_classification_metrics(y_valid, y_valid_preds, threshold=0.3))

```

```

Training: {'Accuracy': 0.9095324833029751, 'Precision': 0.5918572505654687, 'Recall': 0.6345998383185125, 'F1-score': 0.612483745123537
[ 1356, 2355]], 'AUC': 0.9356512676698363}
Validation: {'Accuracy': 0.9043224866440019, 'Precision': 0.5660377358490566, 'Recall': 0.646551724137931, 'F1-score': 0.60362173038229
[ 164, 300]], 'AUC': 0.938864234754544}

```


✓ KNN

```
from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors = 150)
knn.fit(X_train, y_train)
```

```
▼ KNeighborsClassifier
KNeighborsClassifier(n_neighbors=150)
```

```
y_train_predsknn = knn.predict_proba(X_train)[:,:1]
y_valid_predsknn = knn.predict_proba(X_valid)[:,:1]
```

```
print('Training:',calculate_classification_metrics(y_train, y_train_predsknn, threshold=0.3))
print('Validation:',calculate_classification_metrics(y_valid, y_valid_predsknn, threshold=0.3))
```

```
Training: {'Accuracy': 0.8918943533697632, 'Precision': 0.5280059746079163, 'Recall': 0.3810293721368903, 'F1-score': 0.4426357802473,
 [ 2297, 1414]]}, 'AUC': 0.8989852143591607}
Validation: {'Accuracy': 0.8878096163186012, 'Precision': 0.5028571428571429, 'Recall': 0.3793103448275862, 'F1-score': 0.4324324324324
 [ 288, 176]]}, 'AUC': 0.9005314794368005}
```


```
from sklearn.metrics import roc_curve
import matplotlib.pyplot as plt
```

```
thresholds = [0, 0.3, 0.4, 0.5]
plt.figure(figsize=(8, 6))
```

```
for threshold in thresholds:
    # Plot ROC curve for training set
    y_train_pred = (y_train_predsknn >= threshold).astype(int)
    fpr_train, tpr_train, _ = roc_curve(y_train, y_train_pred)
    plt.plot(fpr_train, tpr_train, lw=1, color='blue', alpha=0.5)
    plt.plot(fpr_train, tpr_train, 'o', color='blue', markersize=5) # Add filled dots
    for fpr, tpr in zip(fpr_train, tpr_train):
        plt.text(fpr, tpr, f'{threshold:.1f}', color='blue', fontsize=8, ha='right', va='bottom')
```

```
# Plot ROC curve for validation set
y_valid_pred = (y_valid_predsknn >= threshold).astype(int)
fpr_valid, tpr_valid, _ = roc_curve(y_valid, y_valid_pred)
plt.plot(fpr_valid, tpr_valid, lw=1, color='red', alpha=0.5)
plt.plot(fpr_valid, tpr_valid, 'o', color='red', markersize=5) # Add filled dots
for fpr, tpr in zip(fpr_valid, tpr_valid):
    plt.text(fpr, tpr, f'{threshold:.1f}', color='red', fontsize=8, ha='right', va='bottom')
```

```
plt.plot([0, 1], [0, 1], color='gray', lw=1, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Training and Validation Sets')
plt.legend(['Training Set', 'Validation Set'], loc="lower right")
plt.show()
```



▼ GaussianNB

```
from sklearn.naive_bayes import GaussianNB
```

```
nb = GaussianNB()
nb.fit(X_train, y_train)
```

```
▼ GaussianNB
GaussianNB()
```

```
y_train_predsNb = nb.predict_proba(X_train)[: ,1]
y_valid_predsNb = nb.predict_proba(X_valid)[: ,1]
```

```

from sklearn.metrics import roc_curve
import matplotlib.pyplot as plt


thresholds = [0, 0.3, 0.4, 0.5]
plt.figure(figsize=(8, 6))

for threshold in thresholds:
    # Plot ROC curve for training set
    y_train_pred = (y_train_predsNb >= threshold).astype(int)
    fpr_train, tpr_train, _ = roc_curve(y_train, y_train_pred)
    plt.plot(fpr_train, tpr_train, lw=1, color='blue', alpha=0.5)
    plt.plot(fpr_train, tpr_train, 'o', color='blue', markersize=5) # Add filled dots
    for fpr, tpr in zip(fpr_train, tpr_train):
        plt.text(fpr, tpr, f'{threshold:.1f}', color='blue', fontsize=8, ha='right', va='bottom')

    # Plot ROC curve for validation set
    y_valid_pred = (y_valid_predsNb >= threshold).astype(int)
    fpr_valid, tpr_valid, _ = roc_curve(y_valid, y_valid_pred)
    plt.plot(fpr_valid, tpr_valid, lw=1, color='red', alpha=0.5)
    plt.plot(fpr_valid, tpr_valid, 'o', color='red', markersize=5) # Add filled dots
    for fpr, tpr in zip(fpr_valid, tpr_valid):
        plt.text(fpr, tpr, f'{threshold:.1f}', color='red', fontsize=8, ha='right', va='bottom')

plt.plot([0, 1], [0, 1], color='gray', lw=1, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Training and Validation Sets')
plt.legend(['Training Set', 'Validation Set'], loc="lower right")
plt.show()

```



```

print('Training:', calculate_classification_metrics(y_train, y_train_predsNb, threshold=0.5))
print('Validation:', calculate_classification_metrics(y_valid, y_valid_predsNb, threshold=0.5))

Training: {'Accuracy': 0.7460837887067395, 'Precision': 0.2793322583704828, 'Recall': 0.7935866343303691, 'F1-score': 0.413217342500350
[ 766, 2945]]}, 'AUC': 0.8301500175824722}
Validation: {'Accuracy': 0.7479358912093249, 'Precision': 0.28323262839879154, 'Recall': 0.8081896551724138, 'F1-score': 0.419463087248
[ 89, 375]]}, 'AUC': 0.8331773281052413}

```

```

from sklearn.metrics import roc_auc_score, accuracy_score, precision_score, recall_score, f1_score
def calc_specificity(y_actual, y_pred, thresh):
    # calculates specificity
    return sum((y_pred < thresh) & (y_actual == 0)) / sum(y_actual == 0)

def print_report(y_actual, y_pred, thresh):

    auc = roc_auc_score(y_actual, y_pred)
    accuracy = accuracy_score(y_actual, (y_pred > thresh))
    recall = recall_score(y_actual, (y_pred > thresh))
    precision = precision_score(y_actual, (y_pred > thresh))
    specificity = calc_specificity(y_actual, y_pred, thresh)
    f1 = 2 * (precision * recall) / (precision + recall)

    print('AUC: %.3f' % auc)
    print('accuracy: %.3f' % accuracy)
    print('recall: %.3f' % recall)
    print('precision: %.3f' % precision)
    print('specificity: %.3f' % specificity)
    print('f1: %.3f' % f1)
    print(' ')
    return auc, accuracy, recall, precision, specificity, f1

```

As We Have Done Stratified Split Based on The Percentage of target -As Data Naturally Have 0[negative] instances more therefore accuracy will eventually grow as

1) Accuracy = Total Correct Predicted Label / Total no. of label

As data have dominance in 0 labels most of the labels will be classified as correct predicted which eventually increases accuracy

As we can see it in Logistic Regression 1 class labels have high probability and 0 class label have low probability

if we increase the threshold it will predict more 0 than 1 and if we decrease the threshold to ex: 0.01 then it will predict more 1's than 0

Threshold 0.3 is good

✓ Hence If we do Stratified Split Using Logistic regression will be helpful

Using stratified train-test split on an imbalanced dataset can still be beneficial, especially if you're primarily concerned with evaluating your model's performance accurately. Here's how it can be helpful:

Unbiased Evaluation:

Stratified train-test split ensures that the class distribution in both the training and testing sets resembles that of the original dataset. This prevents the testing set from being skewed towards the majority class, which could lead to overestimating the model's performance, especially if the minority class is of particular interest.

Generalization Performance:

By maintaining the class distribution in the testing set, you get a more realistic estimate of how well your model will perform on unseen data, which is crucial for assessing its generalization performance. This is important regardless of class imbalance because you want your model to perform well on all classes, not just the majority one.

Robustness to Changes:

Ensuring that the class distribution is maintained in the testing set makes your evaluation more robust to changes in the underlying data distribution. If your dataset changes over time (which is common in real-world scenarios), using a stratified approach helps ensure that your evaluation remains consistent and reliable.

✓ Now Lets Do Undersampling

```

from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from collections import Counter
from imblearn.under_sampling import RandomUnderSampler

X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.2, random_state=42)

# Undersampling the majority class
undersampler = RandomUnderSampler(random_state=42)
X_train, y_train = undersampler.fit_resample(X_train, y_train)

# Check the class distribution after undersampling
print("Class distribution after undersampling:", Counter(y_resampled))

    Class distribution after undersampling: Counter({0: 3668, 1: 3668})

# Then, split the test set into validation and test sets (50% each)
X_valid, X_test, y_valid, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42, stratify=y_temp)

# Undersampling the majority class for valid
undersampler = RandomUnderSampler(random_state=42)
X_valid, y_valid = undersampler.fit_resample(X_valid, y_valid)

# Check the class distribution after undersampling
print("Class distribution after undersampling:", Counter(y_valid))

    Class distribution after undersampling: Counter({0: 486, 1: 486})

# Undersampling the majority class for test
undersampler = RandomUnderSampler(random_state=42)
X_test, y_test = undersampler.fit_resample(X_test, y_test)

# Check the class distribution after undersampling
print("Class distribution after undersampling:", Counter(y_test))

    Class distribution after undersampling: Counter({0: 485, 1: 485})

X_train_scaled=scaler.fit_transform(X_train)
X_valid_scaled=scaler.transform(X_valid)
X_test_scaled=scaler.transform(X_test)
X_train=pd.DataFrame(X_train_scaled,columns=X_train.columns)
X_valid=pd.DataFrame(X_valid_scaled,columns=X_valid.columns)
X_test=pd.DataFrame(X_test_scaled,columns=X_test.columns)

```

Each Split Is Now Balanced

✓ Model Selection

```

from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB

```

Logistic Regression

```

lr=LogisticRegression(random_state = 42)
lr.fit(X_train, y_train)

y_train_preds_lr = lr.predict_proba(X_train)[: ,1]
y_valid_preds_lr = lr.predict_proba(X_valid)[: ,1]

from sklearn.metrics import roc_auc_score, accuracy_score, precision_score, recall_score, f1_score

print('Training:',calculate_classification_metrics(y_train, y_train_preds, threshold=0.4))
print('Validation:',calculate_classification_metrics(y_valid, y_valid_preds, threshold=0.4))

    Training: {'Accuracy': 0.8653217011995638, 'Precision': 0.8375314861460957, 'Recall': 0.9064885496183206, 'F1-score': 0.870646766169154
    [ 343, 3325]]}, 'AUC': 0.9339459488707785}

```

```
Validation: {'Accuracy': 0.8446502057613169, 'Precision': 0.8107606679035251, 'Recall': 0.8991769547325102, 'F1-score': 0.8526829268292
[ 49, 437]]}, 'AUC': 0.9152314179749022}
```


```
from sklearn.metrics import roc_curve
import matplotlib.pyplot as plt

thresholds = [0, 0.3, 0.4, 0.5]
plt.figure(figsize=(8, 6))

for threshold in thresholds:
    # Plot ROC curve for training set
    y_train_pred = (y_train_preds1r >= threshold).astype(int)
    fpr_train, tpr_train, _ = roc_curve(y_train, y_train_pred)
    plt.plot(fpr_train, tpr_train, lw=1, color='blue', alpha=0.5)
    plt.plot(fpr_train, tpr_train, 'o', color='blue', markersize=5) # Add filled dots
    for fpr, tpr in zip(fpr_train, tpr_train):
        plt.text(fpr, tpr, f'{threshold:.1f}', color='blue', fontsize=8, ha='right', va='bottom')

    # Plot ROC curve for validation set
    y_valid_pred = (y_valid_preds1r >= threshold).astype(int)
    fpr_valid, tpr_valid, _ = roc_curve(y_valid, y_valid_pred)
    plt.plot(fpr_valid, tpr_valid, lw=1, color='red', alpha=0.5)
    plt.plot(fpr_valid, tpr_valid, 'o', color='red', markersize=5) # Add filled dots
    for fpr, tpr in zip(fpr_valid, tpr_valid):
        plt.text(fpr, tpr, f'{threshold:.1f}', color='red', fontsize=8, ha='right', va='bottom')

plt.plot([0, 1], [0, 1], color='gray', lw=1, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Training and Validation Sets')
plt.legend(['Training Set', 'Validation Set'], loc="lower right")
plt.show()
```



▼ GaussianNb

```
from sklearn.naive_bayes import GaussianNB
```

```
nb = GaussianNB()
nb.fit(X_train, y_train)
```

```
y_train_predsNb = nb.predict_proba(X_train)[:,1]
```

```

y_valid_predsNb = nb.predict_proba(X_valid)[0,1]

from sklearn.metrics import roc_curve
import matplotlib.pyplot as plt

thresholds = [0, 0.3, 0.4, 0.5]
plt.figure(figsize=(8, 6))

for threshold in thresholds:
    # Plot ROC curve for training set
    y_train_pred = (y_train_predsNb >= threshold).astype(int)
    fpr_train, tpr_train, _ = roc_curve(y_train, y_train_pred)
    plt.plot(fpr_train, tpr_train, lw=1, color='blue', alpha=0.5)
    plt.plot(fpr_train, tpr_train, 'o', color='blue', markersize=5) # Add filled dots
    for fpr, tpr in zip(fpr_train, tpr_train):
        plt.text(fpr, tpr, f'{threshold:.1f}', color='blue', fontsize=8, ha='right', va='bottom')

    # Plot ROC curve for validation set
    y_valid_pred = (y_valid_predsNb >= threshold).astype(int)
    fpr_valid, tpr_valid, _ = roc_curve(y_valid, y_valid_pred)
    plt.plot(fpr_valid, tpr_valid, lw=1, color='red', alpha=0.5)
    plt.plot(fpr_valid, tpr_valid, 'o', color='red', markersize=5) # Add filled dots
    for fpr, tpr in zip(fpr_valid, tpr_valid):
        plt.text(fpr, tpr, f'{threshold:.1f}', color='red', fontsize=8, ha='right', va='bottom')

plt.plot([0, 1], [0, 1], color='gray', lw=1, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Training and Validation Sets')
plt.legend(['Training Set', 'Validation Set'], loc="lower right")
plt.show()

```

